

# Bayesian Data Analysis - Assignment 5

October 22, 2017

## 1 Generalized linear model: Bioassay with Metropolis

proposal distribution: (based on  $J_t(\theta^* | \theta^{t-1}) = N(\theta^* | \theta^{t-1}, c^2 \Sigma)$ )

$$N\left(\begin{pmatrix} \alpha \\ \beta \end{pmatrix}, \begin{pmatrix} 3 & 0 \\ 0 & 3 \end{pmatrix}\right)$$

starting point:  $\theta^0 = (\alpha, \beta)^0 = (0, 5)$

number of chains: 10

the number of samples generated from each chain: 1000

the warm-up length: 200

Python code:

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import multivariate_normal
from scipy.stats import norm

# data
x = np.array([-0.86, -0.30, -0.05, 0.73])
n = np.array([5, 5, 5, 5])
y = np.array([0, 1, 3, 5])

# log posterior distribution with parameter a, b
def log_posterior_distribution(a,b):
    ilogit_abx = 1 / (np.exp(-(a + b * x)) + 1)
    log_p = np.log(np.prod(ilogit_abx**y * (1 - ilogit_abx)**(n -
                                                                y)))

    return log_p

chains = 10 # number of chains
iterations = 1000 # number of iterations
starting_point = [0,5] # starting points
a=np.zeros((iterations,chains)) # parameter a
```

```

b=np.zeros((iterations,chains)) # parameter b
a[0,0]=starting_point[0]
b[0,0]=starting_point[1]

# Metropolis algorithm
for i in range(iterations):
    for j in range(chains):
        # sample a set of proposal parameters  $a \sim N(a[i-1,j],3), b \sim N(b[i-1,j],3)$ 
        a_new = norm.rvs(loc=a[i-1,j],scale=np.sqrt(3),size=1)
        b_new = norm.rvs(loc=b[i-1,j],scale=np.sqrt(3),size=1)
        # calculate the new and the previous one log posterior distributions
        log_p_new = log_posterior_distribution(a_new,b_new)
        log_p_prev = log_posterior_distribution(a[i-1,j],b[i-1,j])
        # generate a random number between 0 and 1
        rand = np.random.uniform(0,1)
        # calculate the ratio of the densities
        # do comparision and set parameters (BDA3 P278)
        if min(np.log(1),log_p_new - log_p_prev)>np.log(rand):
            a[i,j]= a_new
            b[i,j]= b_new
        else:
            a[i,j]=a[i-1,j]
            b[i,j]=b[i-1,j]

# (remove the 200 warm-up samples)
a=a[200:]
b=b[200:]

# scatter plot
plt.figure()
plt.scatter(a,b,s=10,edgecolor='black')
plt.xlabel(r'$\alpha$')
plt.ylabel(r'$\beta$')
plt.title("Scatter plot with Metropolis algorithm")

# posterior density
plt.figure()
A = np.linspace(-2, 6, 100)
B = np.linspace(0, 25, 100)
ilogit_abx = 1 / (np.exp(-(A[:,None] + B[:,None,None] * x)) + 1)

p = np.prod(ilogit_abx**y * (1 - ilogit_abx)**(n - y), axis=2)
plt.contourf(p, origin='lower', aspect='auto',
extent=(A[0], A[-1], B[0], B[-1]))
plt.xlim([-2,6])
plt.ylim([0,25])
plt.ylabel(r'$\beta$')
plt.xlabel(r'$\alpha$')

```

```

plt.grid('off')
plt.title('posterior density')

# psrf
def psrf(samples):
    # Calculate means  $\bar{W}$  of the variances
    W = np.mean(np.var(samples,axis=1,ddof=1),axis=0)
    # Calculate variances  $B$  (in fact  $B/n$ ) of the means
    Bpn = np.var(np.mean(samples,axis=1),ddof=1,axis=0)
    B = Bpn*800
    Vh = (800-1)/800*W +Bpn
    R = np.sqrt(Vh/W)
    return R

print("The R of a: {:.6f}".format(psrf(a)))
print("The R of b: {:.6f}".format(psrf(b)))

plt.show()

```

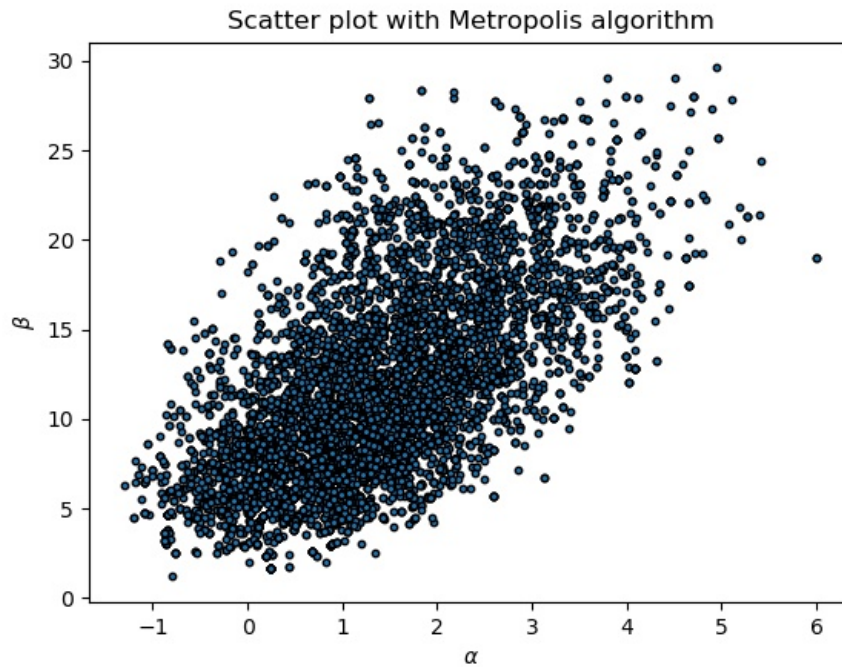


Figure 1: scatter plot with Metropolis algorithm




Figure 2: posterior density

$\hat{R}$  of  $\alpha$ : 1.033702

$\hat{R}$  of  $\beta$ : 1.023332

The scatter plot matches the figure 2.

The  $\hat{R}$  of  $\alpha$  and the  $\hat{R}$  of  $\beta$  are both very close to 1. So we can conclude that the chains have been converged (the samples were from the same distribution).

## 2 Generalized linea rmodel: Bioassay with Stan

bioassay.stan:

```
data{
  int<lower=0> J; // number of doses
  vector[J] x; // values of doses
  int<lower=0> n[J]; // number of animals
```

```

    int <lower=0> y[J]; // number of deaths
  }
parameters{
  real alpha;
  real beta;
}
transformed parameters{
  vector[J] logits;
  logits = alpha + beta*x; // Link function
}
model{
  y~binomial_logit(n,logits);
}

```

R code:

```

library("rstan")
library("ggplot2")
rstan_options(auto_write = TRUE)
options(mc.cores = parallel::detectCores())

biossary_data <- list(J=4,
                      x=c(-0.86,-0.30,-0.05,0.73),
                      n=c(5,5,5,5),
                      y=c(0,1,3,5))
biossary_fit<-stan(file="biossary.stan",data=biossary_data,iter=1000,chains=10)
print(biossary_fit)
biossary_result<-extract(biossary_fit)
p = data.frame(alpha=biossary_result$alpha,beta=biossary_result$beta)
ggplot(p,aes(alpha,beta))+geom_point(shape=21, fill="blue", color="darkred")

```

Output:

Inference for Stan model: biossary.  
 10 chains, each with iter=1000; warmup=500; thin=1;  
 post-warmup draws per chain=500, total post-warmup draws=5000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
alpha	1.27	0.03	1.08	-0.67	0.51	1.19	1.94	3.64	1376	1.01
beta	11.40	0.15	5.62	3.40	7.08	10.39	14.73	24.71	1319	1.01
logits[1]	-8.53	0.11	4.22	-18.50	-10.88	-7.77	-5.40	-2.52	1481	1.01

logits[2]	-2.15	0.03	1.29	-5.26	-2.85	-1.96	-1.24	-0.11	2298	1.00
logits[3]	0.70	0.02	0.93	-1.05	0.09	0.67	1.28	2.67	1614	1.00
logits[4]	9.60	0.14	4.87	2.50	5.80	8.77	12.58	20.95	1238	1.01
lp__	-6.98	0.03	1.08	-9.83	-7.42	-6.64	-6.20	-5.93	1205	1.01

Samples were drawn using NUTS(diag\_e) at Sun Oct 22 10:40:12 2017.  
 For each parameter, `n_eff` is a crude measure of effective sample size,  
 and `Rhat` is the potential scale reduction factor on split chains (at  
 convergence, `Rhat=1`).




Figure 3: scatter plot with Stan

The scatter plot matches the Figure 2.

From the above results, we can see that the  $\hat{R}$ s of  $\alpha$  and  $\beta$  are both 1.01, which is extremely close to 1. So we can conclude that the chains have been converged well.